

Entanglement Swapping in Quantum Repeater Chains

Monte Carlo Simulation with QuantumSavory.jl

Andrea Gennari

A.Y. 2025/2026, Quantum Computing

Outline

Theory

1. Motivation and quantum repeaters
2. Entanglement swapping protocol
3. Noise model: depolarization

Simulation & Results


4. Simulator architecture
5. Ideal case validation
6. Monte Carlo simulation
7. Numerical vs. theoretical comparison
8. Emergent effects and conclusions

Why Quantum Repeaters?

- ▶ Optical fiber: **exponential losses** with distance
- ▶ **No-cloning theorem**: no classical amplification
- ▶ **Quantum repeaters**: hop-by-hop entanglement distribution

Project Goal

Simulate Alice $\leftrightarrow$ Bob entanglement across N repeaters, measuring **fidelity** and **distribution time** vs noise.




The Protocol: Entanglement Swapping

Key idea: a BSM at the repeater teleports entanglement to the remote nodes.

Each repeater R_k (a parallel process):

1. Holds qubits entangled with both neighbours
2. **Local BSM** on its two qubits only
3. Sends the outcome on a **classical channel**
4. Remote node applies the **Pauli correction**

Repeaters swap **independently** until Alice and Bob share $|\Phi^+\rangle$.



Noise Model: Depolarization

Quantum memories **depolarize** while waiting.

Single-qubit depolarizing channel
(QuantumSavory convention):

$$\mathcal{E}(\rho) = (1 - p) \rho + p \frac{I}{2}$$

Noise tied to waiting time:

$$p(\Delta t) = 1 - e^{-\Delta t/\tau}$$

$\tau = -1/\ln(1 - p_w)$, p_w = per-step depolarization rate.

Fidelity: $F = \frac{1}{4}(1 + \langle XX \rangle + \langle ZZ \rangle - \langle YY \rangle)$

Why Time Matters

If $p_s < 1$, links become ready at different times; the end-to-end pair is delivered when the **slowest** link arrives.

$$T = \max_{i=1}^{N+1} T_i, \quad T_i \sim \text{Geom}(p_s)$$

Each repeater swaps **as soon as** it holds both halves (at the max of its two links), so internal qubits are measured early and decohere *less*; only Alice's and Bob's qubits wait until T .

Simulator Architecture

Language: Julia **Framework:**
QuantumSavory.jl



Linear network: Alice + N repeaters + Bob

- ▶ Alice/Bob: 1 slot; each repeater: 2 slots
- ▶ Linear RegisterNet + classical channels

Discrete-event engine (ProtocolZoo)

Parallel @resumable entities on one ConcurrentSim clock:

- ▶ EntanglerProt: heralded generation per link
- ▶ SwapperProt: asynchronous local BSM
- ▶ EntanglementTracker: classical corrections
- ▶ detector: stops at delivery, measures F, T

Ideal Case Validation

Test: no noise ($p_w = 0$), deterministic ($p_s = 1$) $\Rightarrow$ fidelity **exactly 1.0**, delivered at $T = 1$.

N	F	T	Status
1	1.000	1	OK
2	1.000	1	OK
3	1.000	1	OK
5	1.000	1	OK

Swapping preserves $|\Phi^+\rangle$ perfectly; one heralded-generation window costs $T = 1$ (the ideal case is 1, not 0).


What it Validates

- ▶ Bell pair initialization (EntanglerProt)
- ▶ Local BSM + classical-message corrections
- ▶ Asynchronous swap composition ($\forall N$, via Tracker)
- ▶ Fidelity via $\langle XX \rangle, \langle ZZ \rangle, \langle YY \rangle$

Entry point

```
julia -project=. run_ideal.jl
```


MC Simulation: Distribution Time



Parameters: $p_w = 0$, $M = 500$ runs

- ▶ Time scales as $\sim 1/p_s$
(geometric attempts per link)
- ▶ Bottleneck: **slowest link**
 $T = \max_i T_i$
- ▶ More repeaters $\Rightarrow$ more links $\Rightarrow$ larger max
- ▶ Bands: **95% CI** of the mean,
clipped to $T \geq 1$


MC Simulation: Fidelity vs p_s



Parameters: $N = 3$, $M = 500$ runs

- ▶ Low $p_s \rightarrow$ long wait $\rightarrow$ more depolarization $\rightarrow F$ drops
- ▶ $p_w = 0.10$ degrades far more than $p_w = 0.01$
- ▶ At $p_s = 1$: $T = 1$ and no inter-link waiting, so $F \rightarrow 1$
- ▶ Bands: **95% CI** of the mean, clipped to $[0.25, 1]$
- ▶ **Tradeoff**: reliable links reduce both time and degradation


MC Simulation: Fidelity vs N



Parameters: $p_s = 0.5$, $p_w = 0.05$,
 $M = 500$

- ▶ F decreases monotonically with N
- ▶ Two effects: more noisy qubits + larger $T = \max_i T_i$
- ▶ $F \approx 0.92$ ($N=1$) to $F \approx 0.64$ ($N=7$)
- ▶ Asynchronous swaps keep F well above a synchronous bound
- ▶ Bands: **95% CI** of the mean (widen slightly with N)

MC Simulation: Fidelity Heatmap



Parameters: $N = 3$, $M = 500$

2D mean fidelity in (p_s, p_w) space.

- ▶ **Bottom-right** (yellow): high p_s , low $p_w \Rightarrow F \rightarrow 1$
- ▶ **Top-left** (purple): low p_s , high $p_w \Rightarrow F \rightarrow 0.25$ (maximally mixed)
- ▶ Gradient shows **non-linearity**: at low p_s the waiting time amplifies memory noise
- ▶ Right edge: $F = 1$ for every p_w (zero waiting time)

Theoretical Framework

Distribution Time

For $N + 1$ i.i.d. links with $T_i \sim \text{Geom}(p_s)$:

$$\mathbb{E}[T] = \sum_{t=0}^{\infty} \left[1 - (1 - (1 - p_s)^t)^{N+1} \right]$$

Harmonic approximation:

$$\mathbb{E}[T] \approx \frac{H(N+1)}{p_s}, \quad H(n) = \sum_{k=1}^n \frac{1}{k}$$

Werner Model (asynchronous)

Each depolarized qubit survives $(1 - p_w)^\Delta$ over its own wait Δ : a repeater qubit until its swap $\max(T_{i-1}, T_i)$, an end qubit until T .


$$w_i = (1 - p_w)^{\Delta_i^L + \Delta_i^R}, \quad \Delta_i^L, \Delta_i^R \leq T - T_i$$

Ideal BSMs multiply them:

$$w_{\text{tot}} = \prod_{i=1}^{N+1} w_i, \quad F = \frac{1 + 3 w_{\text{tot}}}{4}$$


Per-qubit waits $\Rightarrow$ higher F than synchronous ($\Delta_i = T - T_i$ for all qubits).

Analysis: MC Time vs Theory




- ▶ MC data closely follow **exact curves** (errors $< 5\%$, $M = 1000$)
- ▶ Harmonic approx. $H(N+1)/p_s$ accurate at **low** p_s (geometric $\rightarrow$ exponential limit), overestimates at high p_s
- ▶ Exact formula captures order statistics of i.i.d. geometric RVs
- ▶ Log growth with N : see appendix

Analysis: MC Fidelity vs Werner (async vs synchronous)




Left: the **asynchronous** per-qubit Werner model matches the discrete-event Monte Carlo within statistical error (overlapping 95% CIs) at every p_s and p_w .



Right: the **synchronous** schedule (all BSMs at T , the earlier model) systematically *underestimates* F for $N \geq 2$.




Analysis: Fidelity Scaling with N



Parameters: $p_s = 0.5$, $p_w = 0.05$

- ▶ **Async** Werner tracks MC at every N (within 95% CI)
- ▶ Why: depolarized pairs *are* Werner states; ideal BSMs multiply w exactly
- ▶ **Synchronous** Werner falls away for $N \geq 2$: the gap grows to ~ 0.21 at $N = 7$
- ▶ The two coincide at $N = 1$ (no internal qubit is freed early)

Emergent Effects



Distribution of F : broad, left-skewed mass down to 0.25 (mixed state); the spikes are the *discrete* values of F .

F vs T : clear **negative correlation**; horizontal bands: F is set by the integer waits $T - T_i$, so it is *discrete* and still varies at fixed T .

Distribution of T : right-skewed with a geometric (exponential) tail, outliers wait much longer.

Conclusions

Main Results

- ▶ Built on the **discrete-event** engine: ProtocolZoo entities run in parallel
- ▶ Swapping preserves $|\Phi^+\rangle$ perfectly (**validated** $N = 1 \dots 5$, $T = 1$)
- ▶ F degrades with N , p_w , low p_s (noise + waiting time)
- ▶ **Async** per-qubit Werner matches MC; the synchronous schedule underestimates F for $N \geq 2$

Future Work

- ▶ Entanglement purification to improve F
- ▶ Classical-communication **latency** + cutoff times (CutoffProt)
- ▶ Non-linear topologies (trees, graphs)
- ▶ More realistic noise models (dephasing, losses)

Code

```
run_ideal.jl | run_simulation.jl |  
run_analysis.jl
```


Thank You

Questions?

Appendix: Simulation Parameters

Parameter	Values	Description
N	1, 3, 5 (sim.) / 1–7 (analysis)	Number of repeaters in the chain
p_s	0.1–1.0, step 0.1	Bell pair generation success probability
p_w	0.01, 0.05, 0.10	Per-time-step depolarization parameter
M	500 (sim.) / 1000–2000 (analysis)	Number of Monte Carlo iterations
τ	$-1/\ln(1 - p_w)$	Depolarization time (QuantumSavory)
attempt	1.0	Duration of one heralded-generation window (= 1 step)
delay	10^{-9}	Classical-channel latency (kept negligible)
bands	95% CI	Confidence interval of the mean ($1.96 s/\sqrt{M}$)
Seed	2025	Global RNG seed (reproducible figures)

Appendix: Time Scaling with N



Parameters: $p_s = 0.5$, $p_w = 0$,
 $M = 1000$

- ▶ $\mathbb{E}[T] = \mathbb{E}[\max_{i \leq N+1} T_i]$ grows **logarithmically** in N
- ▶ Harmonic approximation $H(N+1)/p_s$ overestimates slightly at moderate p_s
- ▶ Doubling the chain adds little latency — waiting is dominated by $1/p_s$